



GARCH MODELS IN R


**How much would you lose
in the best of the 5% worst
cases?**

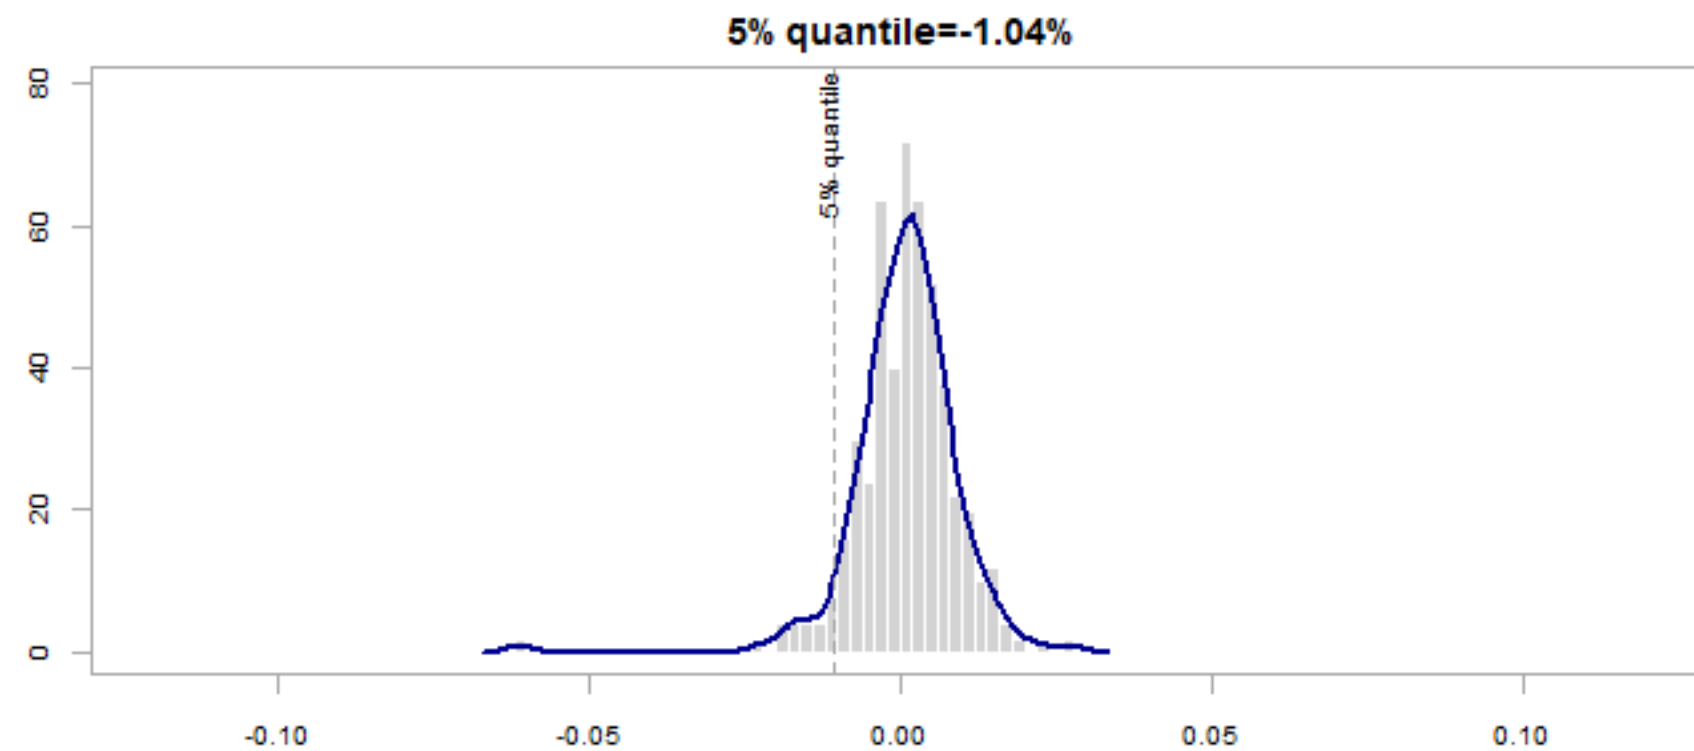

Kris Boudt

Professor of finance and econometrics

Value-at-risk

- A popular measure of downside risk: 5% value-at-risk. The 5% quantile of the return distribution represents the best return in the 5% worst scenarios.





Forward looking approach is needed

- Quantiles of rolling windows of returns are backward looking:
 - ex post question: what has the 5% quantile been for the daily returns over the past year
 - ex ante question: what is the 5% quantile of the predicted distribution of the future return?
- Forward looking risk management uses the predicted quantiles from the GARCH estimation.
- How? Method `quantile()` applied to a `ugarchroll` object.

Workflow to obtain predicted 5% quantiles from ugarchroll

- `ugarchspec()`: Specify which GARCH model you want to use.

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1,0)),  
                        variance.model = list(model = "gjrGARCH"),  
                        distribution.model = "sstd")
```

- `ugarchroll()`: Estimate the GARCH model on rolling estimation samples

```
garchroll <- ugarchroll(garchspec, data = sp500ret, n.start = 2500,  
                      refit.window = "moving", refit.every = 100)
```


- `quantile()`: Compute the predicted quantile

```
garchVaR <- quantile(garchroll, probs = 0.05)
```

- (or any other loss probability that you wish to use: 1% and 2.5% are also popular)

Value-at-risk plot for loss probability 5%

```
actual <- xts(as.data.frame(garchroll)$Realized, time(garchVaR))  
VaRplot(alpha = 0.05, actual = actual, VaR = garchVaR)
```





Exceedance and VaR coverage

A VaR exceedance occurs when the actual return is less than the predicted value-at-risk: $R_t < VaR_t$.

The frequency of VaR exceedances is called the VaR coverage.

```
# Calculation of coverage for S&P 500 returns and 5% probability level  
mean(actual < garchVaR)
```

```
0.05159143
```



VaR coverage and model validation

- Interpretation of coverage for VaR at loss probability α (e.g. 5%):
 - Valid prediction model has a coverage that is close to the probability level α used.
 - If coverage $\gg \alpha$: too many exceedances: the predicted quantile should be more negative. Risk of losing money has been underestimated.
 - If coverage $\ll \alpha$: too few exceedances, the predicted quantile was too negative. Risk of losing money has been overestimated.

Factors that deteriorate the performance

- `distribution.model = "std"` instead of `distribution.model = "sstd"`:

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1,0)),  
                        variance.model = list(model = "gjrGARCH"),  
                        distribution.model = "std")
```

- Rolling estimation and 5% VaR prediction:

```
garchroll <- ugarchroll(garchspec, data = sp500ret, n.start = 2500,  
                      refit.window = "moving", refit.every = 100 )
```

```
garchVaR <- quantile(garchroll, probs = 0.05)
```

```
mean(actual < garchVaR)
```

```
0.05783233
```



Further deterioration

- `variance.model = list(model = "sGARCH")`

instead of

`variance.model = list(model = "gjrgARCH"):`

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1, 0)),  
                        variance.model = list(model = "sGARCH"),  
                        distribution.model = "std")
```

- Rolling estimation and 5% VaR prediction:

```
garchroll <- ugarchroll(garchspec, data = sp500ret, n.start = 2500,  
                      refit.window = "moving", refit.every = 100)
```

```
garchVaR <- quantile(garchroll, probs=0.05)
```

```
mean(actual < garchVaR)
```

```
0.06074475
```



Even further deterioration

- `refit.every = 1000`

instead of

`refit.every = 100:`

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1,0)),  
                        variance.model = list(model = "sGARCH"),  
                        distribution.model = "std")  
  
garchroll <- ugarchroll(garchspec, data = sp500ret, n.start = 2500,  
                      refit.window = "moving", refit.every = 1000)  
  
garchVaR <- quantile(garchroll, probs = 0.05)  
  
mean(actual < garchVaR)  
  
0.06199293
```



GARCH MODELS IN R

**Downside risk means
thinking about predicted
quantiles.**



GARCH MODELS IN R

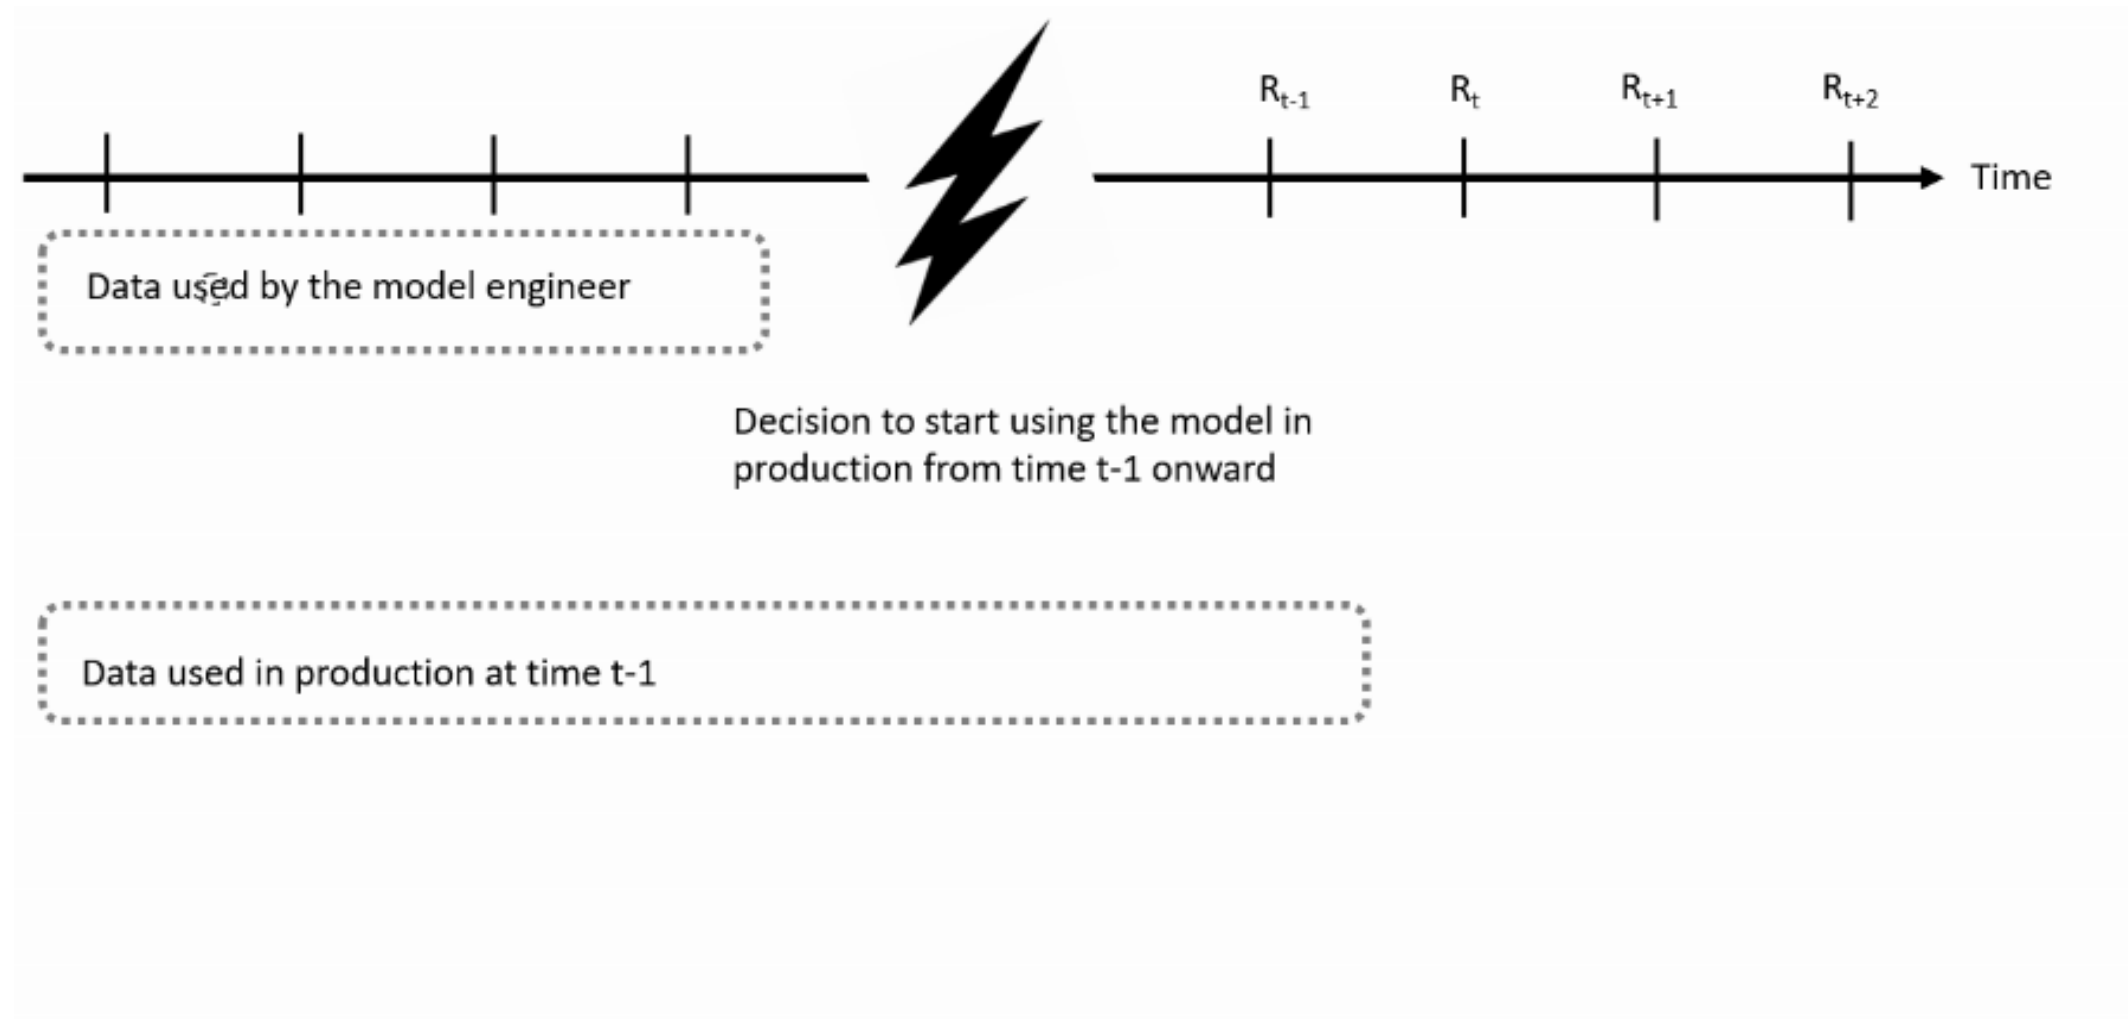
Use the validated GARCH model in production

Kris Boudt

Professor of finance and econometrics



Use in production





New functionality

- Use `ugarchfilter()` for analyzing the recent dynamics in the mean and volatility
- Use `ugarchforecast()` applied to a `ugarchspec` object (instead of `ugarchfit()`) object for making the predictions about the future mean and volatility



Example on MSFT returns

- `msftret`: 1999-2017 daily returns.
- Suppose the model fitting was done using the returns available at year-end 2010.
- You use this model at year-end 2017 to analyze past volatility dynamics and predict future volatility.

Step 1: Defines the final model specification

- Fit the best model using the `msftret` available at year-end 2010:

```
# specify AR(1)-GJR GARCH model with skewed student t distribution
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1,0)),
                        variance.model = list(model = "gjrGARCH"),
                        distribution.model = "sstd")

# estimate the model
garchfit <- ugarchfit(data = msftret["/2010-12"], spec = garchspec)
```

- Define `progarchspec` as the specification to be used in production and use the


instruction `setfixed(progarchspec) <- as.list(coef(garchfit)):`

```
progarchspec <- garchspec
setfixed(progarchspec) <- as.list(coef(garchfit))
```

Step 2: Analysis of past mean and volatility dynamics

- Use the `ugarchfilter()` function:

```
garchfilter <- ugarchfilter(data = msftret, spec = progarchspec)
plot(sigma(garchfilter))
```



Step 3: Make predictions about future returns

```
# Make the predictions for the mean and vol for the next ten days
garchforecast <- ugarchforecast(data = msftret,
                                fitORspec = progarchspec,
                                n.ahead = 10)
```

```
cbind(fitted(garchforecast), sigma(garchforecast))
```

	2017-12-29	2017-12-29
T+1	0.0004781733	0.01124870
T+2	0.0003610470	0.01132550
T+3	0.0003663683	0.01140171
T+4	0.0003661265	0.01147733
T+5	0.0003661375	0.01155238
T+6	0.0003661370	0.01162688
T+7	0.0003661371	0.01170083
T+8	0.0003661371	0.01177424
T+9	0.0003661371	0.01184712
T+10	0.0003661371	0.01191948



Use in simulation

- Instead of applying the complete model to analyze observed returns, you can use it to simulate artificial log-returns:

$$r_t = \log(P_t) - \log(P_{t-1})$$

- Useful to assess the randomness in future returns and the impact on prices, since the future price equals:

$$P_{t+h} = P_t \exp(r_{t+1} + r_{t+2} + \dots + r_{t+h})$$

Step 1: Calibrate the simulation model

- Use the log-returns in the estimation

```
# Compute log returns
msftlogret <- diff(log(MSFTprice)) [(-1)]
```

- Estimate the model and assign model parameters to the simulation model

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1, 0)),
                        variance.model = list(model = "gjrGARCH"),
                        distribution.model = "sstd")


# Estimate the model
garchfit <- ugarchfit(data = msftlogret, spec = garchspec)

# Set that estimated model as the model to be used in the simulation
simgarchspec <- garchspec
setfixed(simgarchspec) <- as.list(coef(garchfit))
```


Step 3: Analysis of simulated returns


- Method `fitted()` provides the simulated returns:

```
simret <- fitted(simgarch)
plot.zoo(simret)
```



Analysis of simulated volatility


```
plot.zoo(sigma(simgarch))
```



Analysis of simulated prices

- Plotting 4 simulations of 10 years of stock prices, with initial price set at 1:

```
simprices <- exp(apply(simret, 2, "cumsum"))  
matplot(simprices, type = "l", lwd = 3)
```





GARCH MODELS IN R

**Time to practice with
`setfixed()`, `ugarchfilter()`,
`ugarchforecast()` and
`ugarchpath()`**



GARCH MODELS IN R

**Model risk is the risk of
using the wrong model**

Kris Boudt

Professor of finance and econometrics



Sources of model risk and solutions

Sources:

- modeling choices
- starting values in the optimization
- outliers in the return series


Solution: Protect yourself through a robust approach

- model-averaging: averaging the predictions of multiple models
- trying several starting values and choosing the one that leads to the highest likelihood
- cleaning the data

Model averaging


- If you cannot choose which model to use, you could estimate them all

```
variance.models <- c("sGARCH", "gjrGARCH")
distribution.models <- c("norm", "std", "std")
c <- 1
for (variance.model in variance.models) {
  for (distribution.model in distribution.models) {
    garchspec <- ugarchspec(mean.model = list(armaOrder = c(0, 0)),
                           variance.model = list(model = variance.model),
                           distribution.model = distribution.model)
    garchfit <- ugarchfit(data = msftret, spec = garchspec)
    if (c==1) {
      msigma <- sigma(garchfit)
    } else {
      msigma <- merge(msigma, sigma(garchfit))
    }
    c <- c + 1
  }
}
```



The average vol prediction

```
avesigma <- xts(rowMeans(msigma), order.by = time(msigma))
```



Robustness to starting values

- GARCH models have many parameters, like

```
coef(garchfit)
```

mu	omega	alpha1	beta1	skew	shape
5.669200e-04	6.281258e-07	7.462984e-02	9.223701e-01	9.436331e-01	6.318621e+00

- Those estimates are the result of a complex optimization of the likelihood function
- Optimization is numeric and iterative: step by step improvement, which can be sensitive to starting values
- `rugarch` has a default approach in getting sensible starting values
- You can specify your own starting values by applying the `setstart()` method to your `ugarchspec()` GARCH model specification

Example with `setstart()` - default starting values

- Estimation with default starting values

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(0,0)),  
                        variance.model = list(model = "sGARCH"),  
                        distribution.model = "sstd")  
  
garchfit <- ugarchfit(data = sp500ret, spec = garchspec)  
  
coef(garchfit)  
  
      mu      omega      alpha1      beta1      skew      shape  
5.669200e-04 6.281258e-07 7.462984e-02 9.223701e-01 9.436331e-01 6.318621e+00  
  
likelihood(garchfit)  
  
24280.33
```

Example with `setstart()` - modified starting values

- Estimation with modified starting values

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(0,0)),  
                        variance.model = list(model = "sGARCH"),  
                        distribution.model = "sstd")  
  
setstart(garchspec) <- list(alpha1 = 0.05, beta1 = 0.9, shape = 8)  
  
garchfit <- ugarchfit(data = sp500ret, spec = garchspec)  
  
coef(garchfit)  
  
      mu      omega      alpha1      beta1      skew      shape  
5.638002e-04 6.303949e-07 7.466503e-02 9.224117e-01 9.438978e-01 6.309185e+00  
  
likelihood(garchfit)  
  
24280.33
```




Cleaning the data

- Avoid that outliers distort the volatility predictions
- How? Through winsorization: reduce the magnitude of the return to an acceptable level using the function `Return.clean()` in the package `PerformanceAnalytics` with `method="boudt"`:

```
# Clean the return series
library(PerformanceAnalytics)
clmsftret <- Return.clean(msftret, method = "boudt")

# Plot them on top of each other
plotret <- plot(msftret, col = "red")
plotret <- addSeries(clmsftret, col = "blue", on = 1)
```




Impact of cleaning on volatility prediction

- Make the volatility predictions using raw and cleaned Microsoft returns

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(1, 0)),  
                        variance.model = list(model = "gjrGARCH"),  
                        distribution.model = "sstd")  
  
garchfit <- ugarchfit(data = msftret, spec = garchspec)  
  
clgarchfit <- ugarchfit(data = clmsftret, spec = garchspec)
```

- Compare them in a time series plot

```
plotvol <- plot(abs(msftret), col = "gray")  
plotvol <- addSeries(sigma(garchfit), col = "red", on = 1)  
plotvol <- addSeries(sigma(clgarchfit), col = "blue", on = 1)  
plotvol
```





GARCH MODELS IN R

**Be a robustnik: it is better
to be roughly right than
exactly wrong**



GARCH MODELS IN R

GARCH volatility leads to time-varying variability of the returns

Kris Boudt

Professor of finance and econometrics



GARCH covariance

- If two asset returns $R_{1,t}$ and $R_{2,t}$ have correlation ρ and time varying volatility $\sigma_{1,t}$ and $\sigma_{2,t}$, then their covariance is:

$$\sigma_{12,t} = \rho \sigma_{1,t} \sigma_{2,t}$$

GARCH covariance estimation in four steps

- Step 1: Use `ugarchfit()` to estimate the GARCH model for each return series.

```
msftgarchfit <- ugarchfit(data = msftret, spec = garchspec)
wmtgarchfit <- ugarchfit(data = wmtret, spec = garchspec)
```

- Step 2: Use `residuals()` to compute the standardized returns.

```
stdmsftret <- residuals(msftgarchfit, standardize = TRUE)
stdwmtret <- residuals(wmtgarchfit, standardize = TRUE)
```

- Step 3: Use `cor()` to estimate ρ as the sample correlation of the standardized returns.


```
msftwmtcor <- as.numeric(cor(stdmsftret, stdwmtret))
msftwmtcor
```

```
0.298795
```

GARCH covariance estimation in four steps

- Step 4: Compute the GARCH covariance by multiplying the estimated correlation and volatilities

```
msftwmtcov <- msftwmtcor * sigma(msftgarchfit) * sigma(wmtgarchfit)
```





Applications of covariance in finance

- Numerous!
- Important case: Optimizing the variance of the portfolio.
- It depends on the:
 - portfolio weights
 - the variance of all the assets
 - the covariance between the asset returns

Application to portfolio optimization

- Variance of portfolio of two assets with weight $w_{1,t}$ invested in asset 1 and $(1 - w_{1,t})$ in asset 2:

$$\sigma_{p,t}^2 = w_{1,t}^2 \sigma_{1,t}^2 + (1 - w_{1,t})^2 \sigma_{2,t}^2 + 2 w_{1,t} (1 - w_{1,t}) \sigma_{12,t}$$

- Many ways to define optimal $w_{1,t}$. One approach is to set $w_{1,t}$ such that the portfolio variance σ_t^2 is minimized.
- First order condition to find minimum variance portfolio

$$\frac{d\sigma_{p,t}^2}{dw_{1,t}} = 2w_{1,t}(\sigma_{1,t}^2 + \sigma_{2,t}^2 - 2\sigma_{12,t}) - 2(\sigma_{2,t}^2 - \sigma_{12,t}) = 0$$




Minimum variance portfolio weights

- Solution:

$$w_{1,t}^* = \frac{\sigma_{2,t}^2 - \sigma_{12,t}}{\sigma_{1,t}^2 + \sigma_{2,t}^2 - 2\sigma_{12,t}}$$

- Calculation in R

```
msftvar <- sigma(msftgarchfit)^2
wmtvar <- sigma(wmtgarchfit)^2
msftwmtcov <- msftwmtcor * sigma(msftgarchfit) * sigma(wmtgarchfit)
msftweight <- (wmtvar - msftwmtcov) / (msftvar + wmtvar - 2 * msftwmtcov)
```



Dynamic beta

- The estimation of a stock's beta: systematic risk of a stock
- Defined as the covariance of the stock return and the market return, divided by the variance of the market returns

$$\beta_t = \frac{\text{covariance between the stock return and the market return}}{\text{variance of the market return}}$$

- Needed to compute the risk premium. The higher it is, the more risky the stock and thus the higher the required rate of return.
- For US stocks, the market return is the return on the S&P 500.

The daily beta of MSFT

- Compute the covariance between MSFT and S&P 500 returns

```
msftsp500cor <- as.numeric(cor(stdmsftret, stdsp500ret))  
msftsp500cov <- msftsp500cor * sigma(msftgarchfit) * sigma(sp500garchfit)
```

- Compute the variance of the S&P 500 returns

```
sp500var <- sigma(sp500garchfit)^2
```


- Compute the beta

```
msftbeta <- msftsp500cov / sp500var
```



msftbeta

1999-01-04 / 2017-12-29





GARCH MODELS IN R

Let's practice!



GARCH MODELS IN R

**Congratulations! You have
learned three languages**

Kris Boudt

Professor of finance and econometrics



Language of GARCH models

- Volatility σ_t at its clusters
- Information set and predictions
- Mean, variance and distribution assumptions
- Leverage effect and the GJR GARCH model
- Skewness, fat tails and the skewed student t distribution
- Model validation using the mean squared error, significance testing, standardized returns and Ljung-Box test
- Applications to value-at-risk, dynamic beta calculation and optimization of financial portfolios



Language of rugarch

- `ugarchspec()`
- `ugarchfit()`
- `ugarchroll()`
- `ugarchforecast()`
- `ugarchfilter()`
- `ugarchpath()`
- **Many useful methods** `sigma()`, `fitted()`, `coef()`, `infocriteria()`,
`likelihood()`, `setfixed()`, `setbounds()`, `quantile()`...



GARCH MODELS IN R

@OptimizeRisk